

## 作业七 MPI 点对点通信实现 Nbody 程序

150\*\*\*\*\* 李师尧

利用 MPI，分别用阻塞式通信、非阻塞通信重新编码实现作业 N-body 计算问题。在同一个多处理机系统上，使用不同的数据规模，对比 MPI 并程序与 pthread 并程序的加速比，分析其加速比差异的原因。

### 1 算法分析

采用如图 1 所示流水并行的方式，设粒子数为  $N$ ，每个进程分配一个 body 和 force 数组，长度为  $N$ ，每个进程负责一小部分的粒子（ $loca \sim locb$ ，数目是  $locbodynum$ ）与其他粒子（ $loca+1 \sim$ ）之间的作用力的计算。每个进程计算作用力之后得到了其他粒子受力的一部分，通过 MPI 点对点通信，将这部分受力加到相应的粒子上（显然，这个过程也可以通过规约操作实现），然后进行时间演化。

计算子集  $p$  中各粒子  $i$  在新时刻  $T'$  的  $\overrightarrow{a(T')_i}$ 、 $\overrightarrow{P(T')_i}$ 、 $\overrightarrow{v(T')_i}$




图 1 N body 问题的流水并行方式

## 2 结果与讨论

如图 2 所示，不同粒子数目下、不同时间步数的性能表现。空心点是 `pthread` 的结果，整体低于 `MPI`，`MPI` 阻塞式通信在小规模时优于非阻塞，但在大规模时，优势不在，根据趋势可知，但规模较大时，阻塞通信使得程序因通信开销大大增加而效率降低，非阻塞通信开始凸显优势。

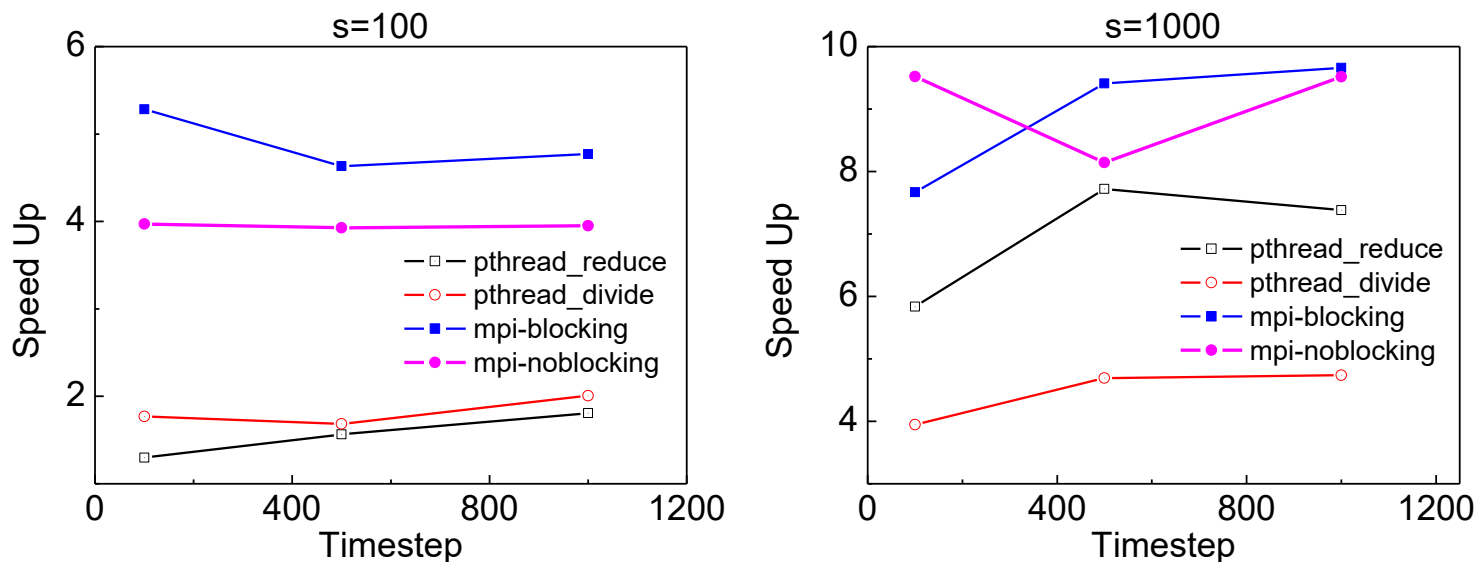


图 2 两种粒子数目、不同时间步数的加速比比较。空心点是 `pthread` 的实现，实心点是 `MPI` 的表现。

```

1  #include <stdio.h>
2  #include <math.h>
3  #include <time.h>
4  #include <string.h>
5  #include <stdlib.h>
6  #include <unistd.h>
7  #include <pthread.h>
8  #include <mpi.h>
9  #define NANO          1000000000
10 #define Max_Thread_Num 256
11 #define REAL          double
12
13 int BodyNum=0;
14 int TimeSteps=0;
15 REAL *body;
16 REAL *force;
17 MPI_Status status;
18 int myid, size;
19 int totalThread;
20 int bcastint[10];
21
22 pthread_cond_t  cond;
23 pthread_mutex_t mtx;
24
25 double serial();
26 double mpi_nbody_blocking();
27 double mpi_nbody_noblocking();
28
29 double pthread_reduce();
30 void *reduce_worker(void *arg);
31 int freeWorker;
32
33 double pthread_divide();
34 void *divide_worker(void *arg);
35 // struct STATUS {
36 //     int lbound;
37 //     int ubound;
38 //     int *task;
39 //     REAL *force;
40 //     pthread_mutex_t  mtx;
41 // }
42 *status;
43
44 int main(int argc, char** argv ) {
45     int i;
46     REAL ser_time, red_time,
47         div_time, mpi_time_blocking, mpi_time_noblocking;

```

```

46     char *pStr;
47
48
49     MPI_Init( &argc, &argv );
50     MPI_Comm_rank( MPI_COMM_WORLD, &myid );
51     MPI_Comm_size( MPI_COMM_WORLD, &size );
52     if (myid==0) {
53         for ( i=1; i<argc; i++ ) {
54             pStr=strstr(argv[i], "-s=");
55             if ( pStr!=NULL) sscanf(pStr, "-s=%d", &BodyNum);
56             pStr=strstr(argv[i], "-t=");
57             if ( pStr!=NULL) sscanf(pStr, "-t=%d", &TimeSteps);
58
59         }
60         if ( BodyNum*TimeSteps==0) {
61             printf("usage: -s=number-of-bodies
62                 -t=number-of-steps\n");
63             return 0;
64         }
65         bcastint[0]=BodyNum;
66         bcastint[1]=TimeSteps;
67     }
68     MPI_Bcast(bcastint,2,MPI_INT,0,MPI_COMM_WORLD);
69     BodyNum=bcastint[0];
70     TimeSteps=bcastint[1];
71
72
73     if (myid==0) {
74         ser_time = serial();
75         printf("serial: %f\n", ser_time);
76     }
77     MPI_Barrier(MPI_COMM_WORLD);
78     mpi_time_blocking=mpi_nbody_blocking();
79     if (myid==0) printf("mpi-blocking: %f
80 speedup=%f\n",mpi_time_blocking,ser_time/mpi_time_blocking);
81     mpi_time_noblocking=mpi_nbody_noblocking();
82     if (myid==0) printf("mpi-noblocking: %f
83 speedup=%f\n",mpi_time_noblocking,ser_time/mpi_time_noblockin
84 g);
85     free(body);
86     free(force);
87     MPI_Finalize();
88 }
89
90 double serial() {
91     REAL fac, fx, fy, fz;
92     REAL dx, dy, dz, sq, dist;

```

```

91     int t, i, j, bi,bj,fi,fj;
92     struct timespec ts,te;
93     double result;
94
95     body = (REAL*)malloc(4*BodyNum*sizeof(REAL));
96     /* Initialize mass and positions in array p to make a
97     test case */
98     for ( i=0; i<BodyNum; i++)    {
99         body[4*i] = 10.05 + i;
100        body[4*i+1] = 30.0*i;
101        body[4*i+2] = 20.0*i;
102        body[4*i+3] = 10.0*i;
103    }
104
105    clock_gettime(CLOCK_REALTIME, &ts);
106
107    force = (REAL*)malloc(3*BodyNum*sizeof(REAL));
108    for ( i=0; i<3*BodyNum; i++)    force[i] = 0;
109
110    t = 0;
111    while ( t<TimeSteps){
112        /* Loop over points calculating force between each
113        pair.*/
114        for ( i=0; i<BodyNum; i++ ) {
115            bi = 4*i;
116            fi = 3*i;
117            for ( j=i+1; j<BodyNum; j++ ) {
118                bj = 4*j;
119                fj = 3*j;
120                /*Calculate force between particle i and j
121                according to Newton's Law*/
122                dx = body[bi+1] - body[bj+1];
123                dy = body[bi+2] - body[bj+2];
124                dz = body[bi+3] - body[bj+3];
125                sq = dx*dx + dy*dy + dz*dz;
126                dist = sqrt(sq);
127                fac = body[bi] * body[bj] / ( dist * sq );
128                fx = fac * dx;
129                fy = fac * dy;
130                fz = fac * dz;
131                /*Add in force and opposite force to particle
132                i and j */
133                force[fi] -= fx;
134                force[fi+1] -= fy;
135                force[fi+2] -= fz;
136                force[fj] += fx;
137                force[fj+1] += fy;
138                force[fj+2] += fz;
139            }
140        }
141        t += TimeSteps;
142    }
143
144    return result;
145 }

```

```

136     }
137     for ( i=0; i<BodyNum; i++ ){
138         bi = 4*i;
139         fi = 3*i;
140         body[bi+1] = body[bi+1] + force[fi] / body[bi];
141         force[fi] = 0;
142         body[bi+2] = body[bi+2] + force[fi+1] / body[bi];
143         force[fi+1] = 0;
144         body[bi+3] = body[bi+3] + force[fi+2] / body[bi];
145         force[fi+2] = 0;
146     }
147     t++;
148 }
149 free(force);
150
151 clock_gettime(CLOCK_REALTIME, &te);
152 result = te.tv_sec - ts.tv_sec +
153         (double)(te.tv_nsec-ts.tv_nsec)/NANO;
154
155 FILE *fResult=fopen("result_ser_nbody.txt", "w");
156 char str[50];
157 for (i=0; i<BodyNum; i++) {
158     sprintf(str, "(%10.4f %10.4f %10.4f %10.4f)\n",
159             body[4*i], body[4*i+1], body[4*i+2], body[4*i+3]);
160     fwrite(str, sizeof(char), strlen(str), fResult);
161 }
162
163
164 double mpi_nbody_blocking() {
165     REAL fac, fx, fy, fz;
166     REAL dx, dy, dz, sq, dist;
167     int t, i, j, bi,bj,fi,fj;
168     struct timespec ts,te;
169     double result;
170     int locbodynum,loca,locb,locnum;
171     REAL *bodysend,*forcesend;
172     if (myid<size-1) {
173         locbodynum=int(BodyNum/size);
174         loca=0+myid*locbodynum;
175         locb=loca+locbodynum-1;
176     }
177     else{
178         locbodynum=BodyNum-int(BodyNum/size)*(size-1);
179         loca=0+(size-1)*int(BodyNum/size);
180         locb=BodyNum-1;
181     }
182     //printf("%d: %d %d %d \n",myid,loca, locb, locbodynum);

```

```

183     body=(REAL*)malloc(4*BodyNum*sizeof(REAL));
184     force=(REAL*)malloc(3*BodyNum*sizeof(REAL));
185     /* Initialize mass and positions in array p to make a
test case */
186     for ( i=0; i<BodyNum; i++)    {
187         bi=4*i;
188         body[bi] = 10.05 + i;
189         body[bi+1] = 30.0*i;
190         body[bi+2] = 20.0*i;
191         body[bi+3] = 10.0*i;
192     }
193
194     clock_gettime(CLOCK_REALTIME, &ts);
195
196     for ( i=0; i<3*locbodynum; i++)    force[loca+i] = 0;
197
198     t = 0;
199     int j1;
200     while ( t<TimeSteps){
201         /* Loop over points calculating force between each
pair.*/
202         for ( i=0; i<locbodynum; i++ ) {
203             bi = loca+4*i;
204             fi = loca+3*i;
205             if (2*i<=BodyNum)
206                 locnum=int(BodyNum/2);
207             else
208                 locnum=BodyNum-int(BodyNum/2);
209             for ( j1=0; j1<locnum; j1++ ) {
210                 j=(j1+i+1)%BodyNum;
211                 bj = 4*j;
212                 fj = 3*j;
213                 /*Calculate force between particle i and j
according to Newton's Law*/
214                 dx = body[bi+1] - body[bj+1];
215                 dy = body[bi+2] - body[bj+2];
216                 dz = body[bi+3] - body[bj+3];
217                 sq = dx*dx + dy*dy + dz*dz;
218                 dist = sqrt(sq);
219                 fac = body[bi] * body[bj] / ( dist * sq );
220                 fx = fac * dx;
221                 fy = fac * dy;
222                 fz = fac * dz;
223                 /*Add in force and opposite force to particle
i and j */
224                 force[fi] -= fx;
225                 force[fi+1] -= fy;
226                 force[fi+2] -= fz;
227                 force[fj] += fx;

```

```

228         force[fj+1] += fy;
229         force[fj+2] += fz;
230     }
231 }
232
233 for (i=0; i<size; i++){
234     if (i==myid) continue;
235
236     MPI_Send(force,3*BodyNum,MPI_DOUBLE,i,myid,MPI_COMM_WORLD);
237 }
238 for (i=0; i<size; i++) {
239     if (i==myid) continue;
240     REAL *temp;
241     temp=(REAL*)malloc(3*BodyNum*sizeof(REAL));
242
243     MPI_Recv(temp,3*BodyNum,MPI_DOUBLE,i,i,MPI_COMM_WORLD,&status);
244     for (j=loca; j<=locb; j++) {
245         fj=3*j;
246         force[fj]+=temp[fj];
247         force[fj+1]+=temp[fj+1];
248         force[fj+2]+=temp[fj+2];
249     }
250
251     delete temp;
252 }
253
254 // REAL *temp;
255 // temp=(REAL*)malloc(3*size*BodyNum*sizeof(REAL));
256 // for (i=0; i<size; i++) {
257 //     if (i==myid) continue;
258 //     MPI_Recv(temp+3*i*BodyNum,3*BodyNum,MPI_DOUBLE,i,i,MPI_COMM_WORLD,&status);
259 // }
260 // for (i=0; i<size; i++) {
261 //     if (i==myid) continue;
262 //     for (j=loca; j<=locb; j++) {
263 //         fj=3*j;
264 //         fi=fj+size*i;
265 //         force[fj]+=temp[fi];
266 //         force[fj+1]+=temp[fi+1];
267 //         force[fj+2]+=temp[fi+2];
268 //     }
269 // }
270 // delete temp;
271
272 for ( i=0; i<locbodynum; i++ ){

```



```

271         bi = 4*i;
272         fi = 3*i;
273         body[bi+1] = body[bi+1] + force[fi] / body[bi];
274         force[fi] = 0;
275         body[bi+2] = body[bi+2] + force[fi+1] / body[bi];
276         force[fi+1] = 0;
277         body[bi+3] = body[bi+3] + force[fi+2] / body[bi];
278         force[fi+2] = 0;
279     }
280
281     for (i=0; i<size; i++){
282         if (i==myid) continue;
283
284         MPI_Send(body+4*loca,4*locbodynum,MPI_DOUBLE,i,my
285         id,MPI_COMM_WORLD);
286     }
287     for (i=0; i<size; i++) {
288         if (i==myid) continue;
289
290         MPI_Recv(body+4*locbodynum*i,4*locbodynum,MPI_DOU
291         BLE,i,i,MPI_COMM_WORLD,&status);
292     }
293
294     t++;
295 }
296 free(force);
297
298 clock_gettime(CLOCK_REALTIME, &te);
299 result = te.tv_sec - ts.tv_sec +
300 (double)(te.tv_nsec-ts.tv_nsec)/NANO;
301
302 FILE *fResult=fopen("result_mpi_nbody.txt", "w");
303 char str[50];
304 for (i=0; i<BodyNum; i++) {
305     sprintf(str, "(%10.4f %10.4f %10.4f %10.4f)\n",
306     body[4*i], body[4*i+1], body[4*i+2], body[4*i+3]);
307     fwrite(str, sizeof(char), strlen(str), fResult);
308 }
309 fclose(fResult);
310 return result;
311 }
312
313 double mpi_nbody_noblocking() {
314     REAL fac, fx, fy, fz;
315     REAL dx, dy, dz, sq, dist;
316     int t, i, j, bi,bj,fi,fj;
317     struct timespec ts,te;
318     double result;

```

```

314 MPI_Request *request;
315 MPI_Status *state;
316 int locbodynum, loca, locb, locnum;
317 REAL *bodysend, *forcesend;
318 if (myid<size-1) {
319     locbodynum=int(BodyNum/size);
320     loca=0+myid*locbodynum;
321     locb=loca+locbodynum-1;
322 }
323 else{
324     locbodynum=BodyNum-int(BodyNum/size)*(size-1);
325     loca=0+(size-1)*int(BodyNum/size);
326     locb=BodyNum-1;
327 }
328 //printf("%d: %d %d %d \n",myid, loca, locb, locbodynum);
329 body=(REAL*)malloc(4*BodyNum*sizeof(REAL));
330 force=(REAL*)malloc(3*BodyNum*sizeof(REAL));
331 /* Initialize mass and positions in array p to make a
test case */
332 for ( i=0; i<BodyNum; i++) {
333     bi=4*i;
334     body[bi] = 10.05 + i;
335     body[bi+1] = 30.0*i;
336     body[bi+2] = 20.0*i;
337     body[bi+3] = 10.0*i;
338 }
339
340 clock_gettime(CLOCK_REALTIME, &ts);
341
342 for ( i=0; i<3*locbodynum; i++) force[loca+i] = 0;
343
344 t = 0;
345 int j1;
346 int req;
347 request=new MPI_Request[size*2];
348 state=new MPI_Status[size*2];
349 while ( t<TimeSteps){
350     /* Loop over points calculating force between each
pair.*/
351     for ( i=0; i<locbodynum; i++ ) {
352         bi = loca+4*i;
353         fi = loca+3*i;
354         if (2*i<=BodyNum)
355             locnum=int(BodyNum/2);
356         else
357             locnum=BodyNum-int(BodyNum/2);
358         for ( j1=0; j1<locnum; j1++ ) {
359             j=(j1+i+1)%BodyNum;
360             bj = 4*j;

```

```

361         fj = 3*j;
362         /*Calculate force between particle i and j
according to Newton's Law*/
363         dx = body[bi+1] - body[bj+1];
364         dy = body[bi+2] - body[bj+2];
365         dz = body[bi+3] - body[bj+3];
366         sq = dx*dx + dy*dy + dz*dz;
367         dist = sqrt(sq);
368         fac = body[bi] * body[bj] / ( dist * sq );
369         fx = fac * dx;
370         fy = fac * dy;
371         fz = fac * dz;
372         /*Add in force and opposite force to particle
i and j */
373         force[fi] -= fx;
374         force[fi+1] -= fy;
375         force[fi+2] -= fz;
376         force[fj] += fx;
377         force[fj+1] += fy;
378         force[fj+2] += fz;
379     }
380 }
381
382 for (i=0; i<size; i++){
383     if (i==myid) continue;
384
385     MPI_Isend(force, 3*BodyNum, MPI_DOUBLE, i, myid, MPI_C
COMM_WORLD, &request[i]);
386
387     REAL *temp;
388     temp=(REAL*)malloc(3*size*BodyNum*sizeof(REAL));
389     req=0;
390     for (i=0; i<size; i++) {
391         if (i==myid) continue;
392
393         MPI_Irecv(temp+3*i*BodyNum, 3*BodyNum, MPI_DOUBLE, i
, i, MPI_COMM_WORLD, &request[size+req]);
394         req++;
395     }
396     MPI_Waitall(size-1, request+size, state+size);
397     //MPI_Barrier(MPI_COMM_WORLD);
398     for (i=0; i<size; i++) {
399         if (i==myid) continue;
400         for (j=loca; j<=locb; j++) {
401             fj=3*j;
402             fi=fj+size*i;
403             force[fj]+=temp[fi];
404             force[fj+1]+=temp[fi+1];
405             force[fj+2]+=temp[fi+2];

```

```

404     }
405 }
406 delete temp;
407
408
409
410 for ( i=0; i<locbodynum; i++ ){
411     bi = 4*i;
412     fi = 3*i;
413     body[bi+1] = body[bi+1] + force[fi] / body[bi];
414     force[fi] = 0;
415     body[bi+2] = body[bi+2] + force[fi+1] / body[bi];
416     force[fi+1] = 0;
417     body[bi+3] = body[bi+3] + force[fi+2] / body[bi];
418     force[fi+2] = 0;
419 }
420
421 for (i=0; i<size; i++){
422     if (i==myid) continue;
423
424     MPI_Isend(body+4*loca,4*locbodynum,MPI_DOUBLE,i,m
425             yid,MPI_COMM_WORLD,&request[i]);
426
427     req=0;
428     for (i=0; i<size; i++) {
429         if (i==myid) continue;
430
431         MPI_Irecv(body+4*locbodynum*i,4*locbodynum,MPI_DO
432                 UBLE,i,i,MPI_COMM_WORLD,&request[size+req]);
433         req++;
434     }
435     MPI_Waitall(size-1,request+size,state+size);
436     //MPI_Barrier(MPI_COMM_WORLD);
437
438     t++;
439 }
440
441 clock_gettime(CLOCK_REALTIME, &te);
442 result = te.tv_sec - ts.tv_sec +
443 (double)(te.tv_nsec-ts.tv_nsec)/NANO;
444
445 FILE *fResult=fopen("result_mpi_nbody.txt", "w");
446 char str[50];
447 for (i=0; i<BodyNum; i++) {
448     sprintf(str, "(%10.4f %10.4f %10.4f %10.4f)\n",
449             body[4*i], body[4*i+1], body[4*i+2], body[4*i+3]);
450     fwrite(str, sizeof(char), strlen(str), fResult);

```

```
447     }  
448     fclose(fResult);  
449     return result;  
450 }
```